

# COMP 3317 Computer Vision

---

## Tutorial 5

TA: Hao Shaozhe

## In This Tutorial

- Reminders & Hints on **Assignment 2**

## ❖ Assignment - 2

### ➤ Topic

- **Digital Image Processing & Feature Extraction** (lecture 2 & 3)
- A **Programming** assignment

### ➤ Submission

- Please submit **a ZIP file** to Moodle by the deadline.
  - ▶ A Python source code file ([assign2.py](#))
  - ▶ A plain text file ([readme.txt](#)) describing the features you have implemented, especially when you have turned in a partially finished implementation.
- Release: **22 February 2024 (Thu)**
- Deadline : **23:59, Mar 6, 2024 (Wed)**
- We **DO NOT accept** any late submission.

## ➤ Download assignment sheets / files



Assignment 2 (Deadline: 23:59, Mar 6, 2024)



Available from **22 February 2024, 4:30 PM** (hidden otherwise)

## Details of Assignment 2:

- Deadline: 23:59, Mar 6, 2024 (Wed)
- Assignment sheet: <[download here](#)>
- Program template: <[download here](#)>
- Sample output: <[download here](#)>
- Sample image: <[download here](#)>

## ➤ If you have any **questions** about Assignment 2, we encourage you to post your questions in the corresponding **discussion forum** on HKU Moodle.



Assignment 2 discussion forum



Available from **22 February 2024, 4:30 PM** (hidden otherwise)

## ❖ Task


- In this assignment, you are going to implement the **functions** for performing
  - 1) color-to-grayscale image conversion
  - 2) corner detection
- To guide your coding, a **partially completed** Python program is provided to you, which provides implementations for
  - parsing the program arguments,
  - loading an input image,
  - plotting the corner detection result,
  - loading/saving the detected corners from/to a file.

## ❖ Task

- In completing this assignment, you only need to modify the following functions:

- 1) `rgb2gray()`
- 2) `smooth1D()`
- 3) `smooth2D()`
- 4) `harris()`

- Please refer to the tutorial notes as well as the comments in the source code for details of these functions.
- You can compare your result against the **sample output** (`corners.lst`) for checking the correctness of your program.



## ❖ Requirements

- Use the formula for the Y-channel of the **YIQ model** in performing the **color-to-grayscale** image conversion.
- Compute  **$I_x$  and  $I_y$**  correctly by finite differences.
- Construct images of  **$I_x^2$ ,  $I_y^2$ , and  $I_x I_y$**  correctly.
- Compute a **proper filter size** for a Gaussian filter based on its **sigma** value.
- Construct a proper **1D Gaussian filter**.
- **Smooth a 2D image** by convolving it with two 1D Gaussian filters.
- Handle the image border using **partial filters** in smoothing.
- Construct an image of the **cornerness function R** correctly.
- Identify potential corners at **local maxima** in the image of the cornerness function R.
- Compute the cornerness value and coordinates of the potential corners up to **sub-pixel accuracy** by **quadratic approximation**.
- Use the **threshold** value to identify strong corners for output.

## ❖ Requirements

- Use the formula for the Y-channel of the YIQ model in performing the color-to-grayscale image conversion.
- Compute  $I_x$  and  $I_y$  correctly by finite differences. ✓ `rgb2gray()`
- Construct images of  $I_x^2$ ,  $I_y^2$ , and  $I_x I_y$  correctly.
- Compute a proper filter size for a Gaussian filter based on its sigma value.
- Construct a proper 1D Gaussian filter.
- Smooth a 2D image by convolving it with two 1D Gaussian filters.
- Handle the image border using partial filters in smoothing.
- Construct an image of the cornerness function R correctly.
- Identify potential corners at local maxima in the image of the cornerness function R.
- Compute the cornerness value and coordinates of the potential corners up to sub-pixel accuracy by quadratic approximation.
- Use the threshold value to identify strong corners for output.



## ➤ `rgb2gray()`


- RGB to YIQ conversion:

$$\begin{bmatrix} Y \\ I \\ Q \end{bmatrix} = \begin{bmatrix} 0.299 & 0.587 & 0.114 \\ 0.596 & -0.275 & -0.321 \\ 0.212 & -0.523 & 0.311 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

- A color image can be easily converted into a monochrome image by taking only the Y component (iluminance) of the above conversion

$$I(i, j) = 0.299 \times R(i, j) + 0.587 \times G(i, j) + 0.114 \times B(i, j)$$

- If you implement `rgb2gray()` correctly, you may get a similar output as follows:



## ❖ Requirements

- Use the formula for the Y-channel of the YIQ model in performing the color-to-grayscale image conversion.
- Compute  $I_x$  and  $I_y$  correctly by finite differences.
- Construct images of  $I_x^2$ ,  $I_y^2$ , and  $I_x I_y$  correctly.
- **Compute a proper filter size for a Gaussian filter based on its sigma value.**
- **Construct a proper 1D Gaussian filter.**
- **Smooth a 2D image by convolving it with two 1D Gaussian filters.** ✓ **smooth1D()**
- **Handle the image border using partial filters in smoothing.** ✓ **smooth2D()**
- Construct an image of the cornerness function R correctly.
- Identify potential corners at local maxima in the image of the cornerness function R.  
Compute the cornerness value and coordinates of the potential corners up to sub-pixel accuracy by quadratic approximation.
- Use the threshold value to identify strong corners for output.

## ➤ `smooth1D()`, `smooth2D()`

- In this assignment, we are going to detect corners in an image with implementing Harris Corner Detection Algorithm.
  - In this algorithm, we are going to make use of a **smoothed** image.
  - In smoothing an image, we need to apply a **Gaussian Filter**.
  - That's why we should complete **`smooth1D()`** and **`smooth2D()`**.
- Perform `smooth2D()` by calling `smooth1D()` twice.
  - In smoothing an image, it involves a 2D convolution. But convolving with a 2D Gaussian Kernel is computation expensive. (lecture 3, pp.31-32)

$$G_{\sigma}(x, y) * I(x, y) = g_{\sigma}(x) * [g_{\sigma}(y) * I(x, y)]$$

- A better approach to perform 2D smoothing is using **two 1D convolutions**. Therefore, forming a 1D Gaussian Kernel is required.

## ❖ smooth1D()

- 1D Gaussian Filter



- Proper filter size – sigma



- Convolve the image with the filter



- Deal with the borders

## ➤ 1D Gaussian Filter

- Gaussian Filter is a class of low-pass filters which is based on the Gaussian Probability Distribution Function.
- Gaussian Distribution is actually a **Normal Distribution**.


$$g(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2} \frac{(x - \mu)^2}{\sigma^2}\right).$$

Here  $\mu$  is the mean value, gives the **location of peak**;

And  $\sigma$  is the variance, controls how wide the peak is.

The bigger  $\sigma$  is, the more we smooth the image.


- ✓ To form a Gaussian filter, we will set  $\mu = 0$ , because we want each pixel to be the one that has the biggest effect on its new, smoothed value.



## ➤ 1D Gaussian Filter

- So, the Gaussian function will be

- 1D Gaussian function ---  $g_{\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}$
- 2D Gaussian function ---  $G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$



- In practice, both the image and the kernel are discrete quantities, and the convolutions are performed as **truncated summations**
- The size of smooth kernel is defined by  **$2n+1$**  (i.e., an odd number).
- For acceptable accuracy, kernels are generally truncated so that the discarded samples are less than  **$1/1000$  of the peak value** (lecture 3, pp. 31)

➤ Proper filter size – sigma

$$\frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}} \leq \left\lceil \frac{1}{\sigma\sqrt{2\pi}} \right\rceil \times \frac{1}{1000}$$




$$x^2 \geq 2\sigma^2 \ln 1000$$



$$|x| \geq \sigma\sqrt{2 \ln 1000}$$

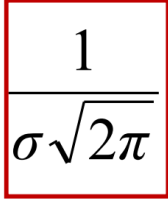
➡ x in this range should be truncated out.





## ➤ 1D Gaussian Filter

- To form a 1D Gaussian filter:

$$g_{\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}$$


Normalization Constant (It is set so that the area under the curve would be 1.)

- ✓ we can **ignore the normalization constant** since we will normalize the output later.

## ➤ 1D Gaussian Filter

- Example

```
import numpy as np
```

```
# find proper filter size  
sigma = 1  
n = int(sigma * (2*np.log(1000))**0.5)  
print(n)
```

```
# form a kernel with the proper size  
x = np.arange(-n, n + 1)  
print(x)
```

```
# create a 1D Gaussian filter  
filter = np.exp((x ** 2) / -2 / (sigma ** 2))  
print(filter)
```

- Output

```
3
```

```
[-3 -2 -1  0  1  2  3]
```

```
[0.011109 0.13533528  
0.60653066 1. 0.60653066  
0.13533528 0.011109 ]
```

## ➤ convolve the image with the filter

- Don't forget to normalize!

```
# Normalize the filter  
filter /= filter.sum()  
print(filter)
```

```
[0.00443305 0.05400558 0.24203623 0.39905028 0.24203623 0.05400558 0.00443305]
```

- Convolve the image with the filter

```
# get the smoothed result  
from scipy.ndimage import convolve1d  
result = convolve1d(img, filter, 1, np.float64, 'constant', 0, 0)
```

## convolve1d

- **Input** : *array\_like*  
The input array.
- **Weights** : *ndarray*  
1-D sequence of numbers.
- **Axis** : *int, optional*  
The axis of *input* along which to calculate. Default is -1.
- **Output** : *array or dtype, optional*  
The array in which to place the output, or the dtype of the returned array. By default an array of the same dtype as input will be created.
- **Cval** : *scalar, optional*  
Value to fill past edges of input if *mode* is 'constant'. Default is 0.0.
- **Origin** : *int, optional*  
Controls the placement of the filter on the input array's pixels. A value of 0 (the default) centers the filter over the pixel, with positive values shifting the filter to the left, and negative ones to the right.

- **Mode** {‘reflect’, ‘constant’, ‘nearest’, ‘mirror’, ‘wrap’}, optional

The *mode* parameter determines how the input array is extended beyond its boundaries. Default is ‘reflect’. Behavior for each valid value is as follows:

- ‘reflect’ (*d c b a | a b c d | d c b a*)

The input is extended by reflecting about the edge of the last pixel. This mode is also sometimes referred to as half-sample symmetric.

- ‘constant’ (*k k k k | a b c d | k k k k*)

The input is extended by filling all values beyond the edge with the same constant value, defined by the *cval* parameter.

- ‘nearest’ (*a a a a | a b c d | d d d d*)

The input is extended by replicating the last pixel.

- ‘mirror’ (*d c b | a b c d | c b a*)

The input is extended by reflecting about the center of the last pixel. This mode is also sometimes referred to as whole-sample symmetric.

- ‘wrap’ (*a b c d | a b c d | a b c d*)

The input is extended by wrapping around to the opposite edge.

- Use `scipy.ndimage` function **`convolve1d`** to conduct 1D conv for smoothing

```
from scipy.ndimage import convolve1d  
res = convolve1d([2, 8, 0, 4, 1, 9, 9, 0], weights=[1, 3])  
print(res)
```

✓ 1.8s

```
[14 24  4 13 12 36 27  0]
```


➤ Before convolving the image with the 1D Gaussian Filter, you have to decide **how to handle the image border**.

✓ In our Assignment 2, **partial filter** is applied

➤ However, we do not have an option for partial filters in **convolve1d()**. Therefore, we need some ways to achieve so.

- You can either choose to deal with the borders separately, or
- Create a matrix whose elements are all ones, and apply the Gaussian Filter to the matrix.
  - each element in this matrix represents a weight.
  - divide the matrix of smoothed image by the matrix of weight.
  - You can skip the normalization for filter if doing so.

## ➤ Explanation





## ❖ `smooth1D()`

- 1D Gaussian Filter



- Proper filter size – sigma



- Convolve the image with the filter





- Deal with the borders

## ❖ smooth2D()

- Call **smooth1D()** twice
  - Convolve each row and column with the 1D Gaussian Filter
  - Use **matrix.T** to transpose a matrix

```
>>> import numpy as np
>>> a = np.array([(1, 2, 3), (4, 5, 6), (7, 8, 9)])
>>> a
array([[1, 2, 3],
       [4, 5, 6],
       [7, 8, 9]])
>>> a.T
array([[1, 4, 7],
       [2, 5, 8],
       [3, 6, 9]])
```

- If you implement `smooth1D()` and `smooth2D()` correctly, you may get a similar output as follows (if applied on grayscale image):



## ❖ Requirements

- Use the formula for the Y-channel of the YIQ model in performing the color-to-grayscale image conversion.
- **Compute  $I_x$  and  $I_y$  correctly by finite differences.**
- **Construct images of  $I_x^2$ ,  $I_y^2$ , and  $I_x I_y$  correctly.**
- Compute a proper filter size for a Gaussian filter based on its sigma value.
- Construct a proper 1D Gaussian filter.
- Smooth a 2D image by convolving it with two 1D Gaussian filters. ✓ **`harris()`**
- Handle the image border using partial filters in smoothing.
- **Construct an image of the cornerness function R correctly.**
- **Identify potential corners at local maxima in the image of the cornerness function R.**
- **Compute the cornerness value and coordinates of the potential corners up to sub-pixel accuracy by quadratic approximation.**
- **Use the threshold value to identify strong corners for output.**

# Corner Detection

- Summary of Harris corner detection algorithm:
  1. Compute  $I_x$  and  $I_y$  at each pixel  $I(x, y)$
  2. Form the images of  $I_x^2$ ,  $I_y^2$  and  $I_x I_y$  respectively
  3. Smooth the images of squared image derivatives
  4. Form an image of the corneriness function  $R$  using the smoothed images of squared derivatives (i.e.,  $\langle I_x^2 \rangle$ ,  $\langle I_y^2 \rangle$  and  $\langle I_x I_y \rangle$ )
  5. Locate local maxima in the image of  $R$  as corners
  6. Compute the coordinates of the corners up to sub-pixel accuracy by quadratic approximation using values in the neighborhood
  7. Threshold the corners so that only those with a value of  $R$  above a certain value are retained

- Compute  $I_x$  and  $I_y$  correctly by finite differences.
  - Compute the **Horizontal** ( $I_x$ ) and **Vertical** ( $I_y$ ) Derivatives


- Construct images of  $I_x^2$ ,  $I_y^2$ , and  $I_x I_y$  correctly.

- $I_x^2 = I_x \times I_x$
- $I_y^2 = I_y \times I_y$
- $I_x I_y = I_x \times I_y$

➡

$$\mathbf{A} = \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Smooth  $I_x^2$ ,  $I_y^2$ , and  $I_x I_y$



- Use numpy function `np.gradient` to calculate  $I_x I_y$

```
import numpy as np
```

```
arr = np.array([[6, 3, 5, 1],  
               [3, 8, 6, 3],  
               [7, 3, 3, 4]])
```

```
grad = np.gradient(arr)
```

```
grad0 = np.gradient(arr, axis=0)
```

```
grad1 = np.gradient(arr, axis=1)
```

```
print(arr)
```

```
print(grad)
```

```
print(grad0)
```

```
print(grad1)
```

✓ 0.5s

```
[[6 3 5 1]
```

```
 [3 8 6 3]
```

```
 [7 3 3 4]]
```

```
[array([-3. ,  5. ,  1. ,  2. ],
```

```
       [ 0.5,  0. , -1. ,  1.5],
```

```
       [ 4. , -5. , -3. ,  1. ]]), array([-3. , -0.5, -1. , -4. ],
```

```
       [ 5. ,  1.5, -2.5, -3. ],
```

```
       [-4. , -2. ,  0.5,  1. ]]])
```

```
[-3.   5.   1.   2. ]
```

```
[ 0.5  0.  -1.   1.5]
```

```
[ 4.  -5.  -3.   1. ]]
```

```
[-3.  -0.5 -1.  -4. ]
```

```
[ 5.   1.5 -2.5 -3. ]
```

```
[-4.  -2.   0.5  1. ]]
```

➤ **Cornerness function R** [lecture 3, pp. 27]

- In Harris corner detection algorithm, corners are marked at points where the quantity  $R = \lambda_1 \lambda_2 - \kappa(\lambda_1 + \lambda_2)^2$  exceeds some threshold, here  $\kappa$  is a parameter set to 0.04 as suggested by Harris
- Note that  $\lambda_1 \lambda_2 = \det(\mathbf{A})$  and  $\lambda_1 + \lambda_2 = \text{trace}(\mathbf{A})$

$$\mathbf{A} = \begin{bmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{bmatrix} \quad \det(\mathbf{A}) = a_{00}a_{11} - a_{01}a_{10} \quad \text{trace}(\mathbf{A}) = a_{00} + a_{11}$$

and hence direct computation of the eigenvalues is not necessary

- Corners are defined as **local maxima** of the cornerness function  $R$ , and sub-pixel accuracy is achieved through a **quadratic approximation** of the neighborhood of the local maxima (using **4-neighbours** and the center pixel):


$$f(x,y) = ax^2 + by^2 + cx + dy + e$$



## ➤ Calculate cornerness function R

$$\mathbf{A} = \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \quad \rightarrow \quad \begin{cases} \det(\mathbf{A}) = I_x^2 I_y^2 - (I_x I_y)^2 \\ \text{trace}(\mathbf{A}) = I_x^2 + I_y^2 \end{cases}$$

$$\begin{aligned} \rightarrow R &= \lambda_1 \lambda_2 - \kappa (\lambda_1 + \lambda_2)^2 \\ &= \det(\mathbf{A}) - \kappa (\text{trace}(\mathbf{A}))^2 & \kappa &= 0.04 \\ &= (I_x^2 I_y^2 - (I_x I_y)^2) - 0.04 (I_x^2 + I_y^2)^2 \end{aligned}$$



## ➤ Find the local maxima

- ✓ Perform **non-maximal suppression** by considering **8-neighbors**

2	2	4	7	5
1	3	4	5	8
1	2	2	4	3
3	8	7	9	6
1	7	8	3	6

**R**




2	2	4	7	5
1	3	4	5	8
1	2	2	4	3
3	8	7	9	6
1	7	8	3	6

**R**


- `smooth1D()`

⇒ Consider boundary. (partial filter)



- Find local maxima

⇒ Do not Consider boundary




# 1D Edge Detection

- Having obtained the derivative  $S'(x)$ , interpolation can be used to locate any maxima or minima to **sub-pixel accuracy**

- Approximate the function locally by  $f(x) = ax^2 + bx + c$
- Without loss of generality, let the sample maximum and its immediate neighbors have coordinates  $x = 0, -1$  and  $1$  respectively
- This gives

$$\begin{cases} f(-1) = a - b + c \\ f(0) = c \\ f(1) = a + b + c \end{cases} \Rightarrow \begin{cases} a = \frac{f(1) + f(-1) - 2f(0)}{2} \\ b = \frac{f(1) - f(-1)}{2} \\ c = f(0) \end{cases}$$



- Locate the maximum/minimum by solving  $f'(x) = 2ax + b = 0$  which gives


$$x_e = -\frac{b}{2a} = -\frac{f(1) - f(-1)}{2[f(1) + f(-1) - 2f(0)]}$$

- Finally, an edge is marked at each maximum or minimum whose magnitude exceeds some thresholds

## ➤ 2D Quadratic approximation

- Perform quadratic approximation to local corner up to **sub-pixel accuracy** (using 4-neighbours and the center pixel)

$$f(x,y)=ax^2+by^2+cx+dy+e$$



## ➤ 2D Quadratic approximation


$$f(x,y)=ax^2+by^2+cx+dy+e$$



$$\begin{aligned}f(0,0) &= e \\f(-1,0) &= a - c + e \\f(1,0) &= a + c + e \\f(0,-1) &= b - d + e \\f(0,1) &= b + d + e\end{aligned}$$



$$\begin{aligned}a &= \frac{f(-1,0) + f(1,0) - 2f(0,0)}{2} \\b &= \frac{f(0,-1) + f(0,1) - 2f(0,0)}{2} \\c &= \frac{f(1,0) - f(-1,0)}{2} \\d &= \frac{f(0,1) - f(0,-1)}{2} \\e &= f(0,0)\end{aligned}$$


$$\left\{ \begin{aligned}x &= -\frac{c}{2a} \\y &= -\frac{d}{2b}\end{aligned} \right.$$

## ➤ Use the threshold value to identify strong corners for output.

- Cornerness > Threshold.  $\Rightarrow$  consider it as a corner

➤ Follow the hints given in program template

```
# TODO: compute  $I_x$  &  $I_y$ 

# TODO: compute  $I_{x^2}$ ,  $I_{y^2}$  and  $I_{xy}$ 

# TODO: smooth the squared derivatives

# TODO: compute corneness function  $R$ 

# TODO: mark local maxima as corner candidates;
#   perform quadratic approximation to local corners upto sub-pixel
accuracy

# TODO: perform thresholding and discard weak corners
```